

Terraform-Basics

- Infrastructure as Code:
 - Introduction: Further Read: [Infrastructure as Code | IBM](#)
 - Concept where the infrastructure requirements and dependencies are written in a file (in a predefined language).
 - The file is maintained as the definition for infrastructure
 - Optionally stored & maintained in a source code repository[best practise].
 - The definition is used by a tool to provision infrastructure.
 - Advantage:
 - Automated deployment
 - Consistent environments
 - Repeatable process
 - Reusable components
 - Documented architecture

- Generic Terms:

Idempotent	denoting an element of a set which is unchanged in value when multiplied or otherwise operated on by itself.
Mutable	Something that keeps changing or fluctuating.
Declarative	denoting high-level programming languages which can be used to solve problems without requiring the programmer to specify an exact procedure to be followed
Procedural	a type of computer programming language that specifies a series of well-structured steps and procedures within its programming context to compose a program.

- Terraform:

- Introduction:
 - Built by Hashicorp. It is an Infrastructure provisioning tool.
 - Example of implementation of Infrastructure as Code.
 - Three Editions:

Terraform CLI	Open Source - Free to use. Single code binary file
Terraform Cloud	Commercial public as a service offering. Hosted @ https://app.terraform.io .
Terraform Enterprise	Terraform Cloud implementation in Private instance

- Note: Terraform documentation is well written and self-explanatory to understand terraform.
- Links:
 - Generic - [Terraform by HashiCorp](#)
 - Registry - [Terraform Registry](#)

- Components

Executable	• Binary that performs the terraform activity
State file	• A source of truth file for terraform. Contains the details of the infrastructure that needs to work on. • Has extension of .tfstate . Is in JSON format. • Changes will be updated by terraform. Not to be updated externally(corrupts file). • Types: <ul style="list-style-type: none">• Local - on same host where terraform is running• Remote - on a different host (can be public cloud, terraform cloud etc)
Plugins	• Code modules that serve as Extensions to terraform.

	<ul style="list-style-type: none"> One such plugin is called Provider, which defines which component to be worked on . Browse Providers Terraform Registry
Definition files (terraform files)	<ul style="list-style-type: none"> Written in HCL (Hashicorp config language) to define what infrastructure is required.

- Typical type of files:

*.tf	Actual code in HCL format
*.var	variable definitions
*.tstate	State file
*.tfvars	Variable value assignments

- How it works (core workflow & commands):

Terraform Init	Initializes the directory structure to be used by Terraform. Command: \$ terraform init
Terraform validate	Validates if the configuration files are correct and there are no errors.
Terraform Plan	<p>The first part of the execution [with option to save the plan]. Steps Involved:</p> <ul style="list-style-type: none"> Uses the definition to connect to the provider collect the necessary information Compare with state file Identify dependencies list out actions that will be performed by terraform to get the infrastructure to the desired state. Lists out the changes that will be performed <ul style="list-style-type: none"> Will be denoted by <ul style="list-style-type: none"> (+) add (colour coded as green) (-) remove (colour coded as red) (-/+) update - usually would mean remove and then add. (colour coded as yellow) <p>Command : \$ terraform plan [options] terraform file</p> <p>Note:</p> <ul style="list-style-type: none"> If run with -destroy switch instead of creating the resources, will destroy the resources. -var <variable>=<value> is another way to pass variables to tf files. -out <filename> will store the plan in an external file.
Terraform Apply	<ul style="list-style-type: none"> Apply the terraform plan (i.e. make changes to the infrastructure) Can be from the last plan generated or from a plan file (stored as *.tfplan) Will make changes to the tfstate file

- Important :**

- Changes made outside of terraform will not be tracked.
- If there is a change to an infrastructure that is provisioned by terraform outside of terraform. Upon execution, terraform will see it as a variation from desired state and will re-apply per the terraform definition.
- If there are multiple terraform configuration files, terraform will stitch them all together and then generate the plan. So, you do not have to link each file/ provide any dependencies. For coding simplicity, you can segregate each file based on purpose & terraform will pick them all together.
- Another important feature is the execution sequence, terraform does not execute blocks completely based on the sequence they are written. While it does start top-down, it also

checks for dependencies and works on the resources automatically based on those dependencies.

- This dependencies can also be configured manually.

- Hashicorp Configuration Language: [Overview - Configuration Language - Terraform by HashiCorp](#)

- Generic Block format:

- Basic block

```
block_type label_one label_two {  
  key = value  
  embedded_block {  
    key = value  
  }  
}
```

- Recognised datatypes:

string	- strings Example: string = "blah blah blah"
Number	- numeric values (including decimal) Example: number = 5
Boolean	- true / false Example: bool = true
List	- For storing a list of values of same data type. Example: list = ["abc","def"] Referenced as list[1] (starts with 0)
Map	- Key - Value pair (for any recognised data type) Example: map = {name = "abc", age = 4, school=true} Referenced as map["name"] . Key is the index input.

- Variable

- Purpose:

- To store dynamic content

- Definitions:

- variable <variable name>{}

- Variable <variable name> { value assignment}

- Example:

- variable "aws_secret_key" {}

- variable "aws_region" { default = "us-east-1" }

- Example:

- ◆ #Specify default variable and type

- variable "environment_name" {

- type = string

- default = "development"

- }

- ◆ #Specify variable in file

- environment_name = "uat"

- ◆ #Specify variable in-line

- terraform plan -var 'environment_name=production' <-passing
variable during plan

- ◆ #Create variable map

- variable "cidr" {

- type = map(string)

- default = {

```

development = "10.0.0.0/16"
uat = "10.1.0.0/16"
production = "10.2.0.0/16"
}
}
#Use map based on environment
cidr_block = lookup(var.cidr, var.environment_name)

```

- Usage:
 - var.<variable name>
 - Example: region = "var.aws_region"
- Interpolation can be done using regular operations, but variable needs to be enclosed in \${}
 - myimage_name = "myproj-\${var.env_type}"
- Expressions, Functions & Loops:
 - HCL has provisions to work with variables and different data types.
 - Functions: =[Functions - Configuration Language - Terraform by HashiCorp](#)
 - Expressions - [Expressions - Configuration Language - Terraform by HashiCorp](#)
 - Note: HCL is not a full blown program language, so while there are loops & conditions, they are minimalistic and with constraints.
 - For & IF:
 - For Expression:
 - ◆ Usage:
 - [for <var> in <list/tuple/map/object>: [actions - with variable referenced as \${var}]
 - ◆ Example [for s in var.list : upper(s)]
 - IF Expression :
 - ◆ Usage: <CONDITION> ? <TRUE VAL> : <FALSE VAL>
 - ◆ Example:
 - var.a != "" ? var.a : "default-a"
- Locals: - [Simplify Terraform Configuration with Locals | Terraform - HashiCorp Learn](#)
 - Purpose :
 - To represent repeated values in a config file. Something like a constant variable definition.
 - Definition :
 - local {
 variable1 = value
 variable2 = value
 }
 - Referenced using local.<variable>
 - Example:
 - locals {
 required_tags = {
 project = var.project_name,
 environment = var.environment
 }
 tags = merge(var.resource_tags, local.required_tags)
 }
- Provider:
 - Purpose:
 - Defines what type of infrastructure & where it is being provisioned.
 - Mandatory. This block establishes connection to the target.
 - Establishing connection
 - Definition:
 - terraform {
 required_providers {
 <provider name> {

```

source = <path for provider plugin information in terraform
registry>
Version = <version of provider plugin>
[options]
} }

```

- Example:

```

terraform {
  required_providers {
    aws = {
      source = "hashicorp/aws"
      version = "3.52.0"
    }
  }
}

```

- Provider Block:

- Purpose:
 - ◆ Provides configuration options for the provider connection.
- Definition:

```

provider <provider name> {
  provider specific parameters / configuration options
}

```
- Example:

```

provider "aws" {
  access_key = "var.access_key"
  secret_key = "var.secret_key"
  region = "var.aws_region"
}

```

- Data Block:

- Purpose:
 - Store data collected from the target
 - Options and module names are specific to the provider, please check the provider documentation .
 - ◆ Data source to use, need to be published by the provider.
- Definition:

```

data <module name> <variable name> {
  [options]
}

```
- Example: - [aws_ami | Data Sources | hashicorp/aws | Terraform Registry](#)

```

data "aws_ami" "alx" {
  most_recent = true
  owners = ["amazon"]
  filters {}
}

```
- Value is referenced as data.<data sourcename>.<data returned> (chck doc on data type of returned value, as how it is referenced will vary based on that)
Example: data.aws_ami.architecture

- Resource Block: [Resources Overview - Configuration Language - Terraform by HashiCorp](#)

- Purpose:
 - Reference a resource for working with. For data collection, use the data resource in data block.
 - ◆ Example (for same type - AWS AMI):
 - ◇ Data source [aws_ami | Data Sources | hashicorp/aws | Terraform Registry](#)
 - ◇ Resource - [aws_ami | Resources | hashicorp/aws | Terraform Registry](#)
 - Options and resource names are specific to the provider, please check the provider documentation .
 - ◆ Resource to be used need to be published by the provider

- Definition


```
Resource <resource_name> <variable name> {
  [options]
}
```
- Example: [aws_instance | Resources | hashicorp/aws | Terraform Registry](#)

```
resource "aws_instance" "ex"{
  ami = "data.aws_ami.alx.id"
  instance_type = "t2.micro"
}
```
- Additional resource manipulation can be done using the documented meta data
 - Count (for total no of resources needed)
 - ◆ count.index - represents the running count number (max = count)
 - depends_on = to manually specify inter resource dependencies.
Example:


```
resource "aws_instance" "myservers" {
  count = 2
  tags {
    Name = "customer-${count.index}"
  }
  depends_on = [aws_iam_role_policy.allow_s3]
}
```
 - For_each - Resource creations based on conditions. "Each" variable will have the content per iteration.
Example: (if two S3 buckets to be created with name DEV_MyProj & PROD_MyProj). First iteration creates DEV_MyProj with public_read ACL and then creates PROD_MyProj with private ACL .


```
resource "aws_s3_bucket" "mybuckets" {
  for_each = {
    DEV= "public-read"
    PROD = "private"
  }
  bucket = "${each.key}_MyProj"
  acl = each.value
}
```
- Output Block:
 - Purpose:
 - If there is a need to display a result after the plan execution.
 - Like a return value.
 - Definition:


```
output <variable name> {
  value = <value to store>}
```
 - Example :


```
output "aws_public_ip"
{
  value = "aws_instance.ex.public_dns "
```
- Provisioners:
 - Further read - [Provisioners - Terraform by HashiCorp](#)
 - Purpose:
 - An alternate method to perform operations or executions.
 - To be used as a last resort only (as these do not maintain state and hard for tracking)
 - Written within the resource block & usually used to perform some task after the resource is created or destroyed
 - Commonly used provisioners:

File	- File Operations , like copy etc.
------	------------------------------------

Local Exec	- Runs commands/ scripts on local machine
Remote Exec	- Runs commands on target machine

- Definition:

```

provisioner <provisioner type> {
  [options]
}

```
- Example:

```

resource "aws_instance" "web" {
  # ...

  provisioner "file" {
    source     = "script.sh"
    destination = "/tmp/script.sh"
  }

  provisioner "remote-exec" {
    inline = [
      "chmod +x /tmp/script.sh",
      "/tmp/script.sh args",
    ]
  }
}

```

- **Workspaces:** [State: Workspaces - Terraform by HashiCorp](#)

- Purpose - A Concept in terraform where the terraform statefile is separated while keeping the other config files same.
 - Common scenario is if we want to use the same configuration for different environments (like DEV/UAT/PROD), you would create different state files for each environment. In which case you would create 3 workspaces for each environment.
 - You would create the config files to be dynamic so that execution is based on which workspace it is working in. Typically referencing it using `${terraform.workspace}` to identify current workspace
 - Based on workspace, load different variables to which are environment specific
- Commands:

```

terraform workspace new <workspace name>
Terraform workspace select <workspace name> -> to go to respective workspace.

```

- **Module:**

- Two types:
 - Provided by 3rd Party (in terraform registry)
 - User defined
- Purpose: reusable components of code. Think functions or modules in other programs, that take some parameters, do the work and return some values.
- Creating a module:
 - Create a directory and put the necessary code in that directory.
 - If you want a return value, make sure you provide it in the output block.
- Using a module :

```

module <module name> {
  source = <path of the code files> # can be local, consul, github or other locations
}

```

Output is referenced using `module.<resourcename>.<output variable>`